

COL761: Assignment 1 - Part 1

Abhay Sharma
2012CS50272

Akash Dabaniya
2021CS10111

Mridul Singh
2018CS50412

February 4, 2026

1 Task 1: Comparison of Apriori and FP-Tree Algorithms

Introduction

An empirical comparison between two prominent Frequent Itemset Mining (FIM) algorithms: **Apriori** and **FP-Growth** (Frequent Pattern Growth). The objective is to analyze their runtime performance across varying support thresholds (5%, 10%, 25%, 50%, and 90%).

Dataset

The experiments utilize the `webdocs.dat` dataset.

- **Characteristics:** The dataset is *dense* and contains very long transactions, unlike typical sparse market-basket data.

Algorithm Overview

Apriori Algorithm

The Apriori algorithm uses a *generate-and-test* approach, originally proposed by Agrawal and Srikant [1]. It relies on the **Apriori Property** (or Downward Closure Property), which states that:

All non-empty subsets of a frequent itemset must also be frequent.

How it works:

1. It works iteratively (level-wise). In the k -th pass, it generates candidate itemsets of length k (C_k) using frequent itemsets from the previous pass (L_{k-1}).
2. It scans the entire database to count the support for these candidates.
3. An itemset X is considered frequent if its support meets the threshold:

$$Support(X) = \frac{\text{count}(X)}{|D|} \geq \text{min_sup}$$

Pros & Cons:

- + **Pros:** Simple to implement and easy to parallelize.
- **Cons:** It requires multiple scans of the database. Crucially, on dense datasets like `webdocs`, it suffers from **candidate explosion**, generating huge numbers of candidates that turn out to be infrequent, wasting significant processing time.

1.0.1 FP-Growth Algorithm

The FP-Growth algorithm adopts a *divide-and-conquer* strategy using a tree structure, avoiding candidate generation entirely, as introduced by Han et al. [2].

How it works:

1. **Compression:** It performs two scans of the database. The first scan counts item frequencies. The second scan compresses the database into a compact structure called the **FP-Tree** (Frequent Pattern Tree) in memory.
2. **Mining:** It mines the tree by recursively building "conditional FP-trees." For a specific item, it looks only at the paths in the tree ending with that item to find frequent prefixes.

Pros & Cons:

- + **Pros:** Extremely efficient because it avoids the cost of generating candidates. It only scans the database twice, regardless of pattern length.
- **Cons:** The FP-Tree is stored in main memory (RAM). For extremely large databases, the tree might become too large to fit.

Experimental Results

The algorithms were executed at support thresholds of 5%, 10%, 25%, 50%, and 90%. The runtimes were recorded and plotted (see Figure 1).

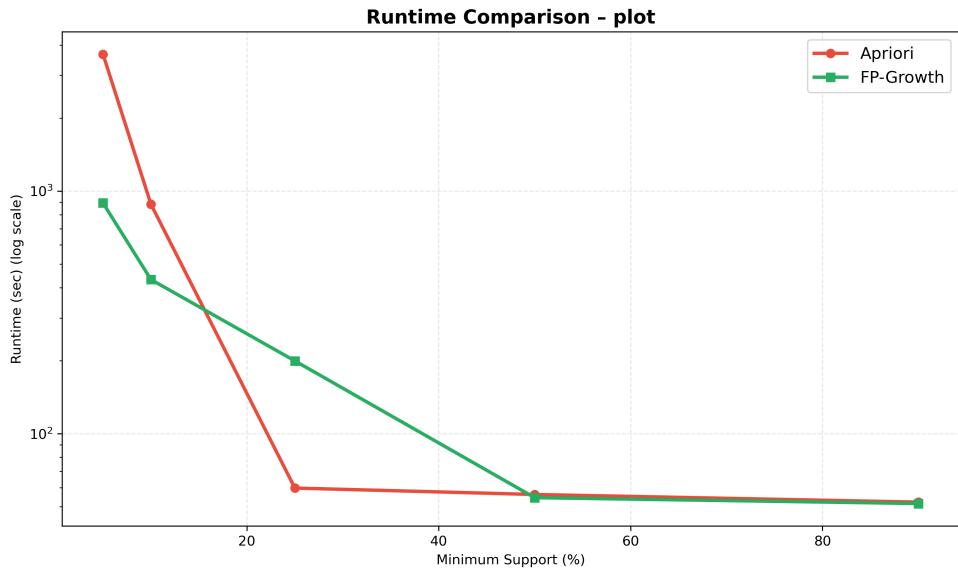


Figure 1: Runtime Comparison of Apriori vs. FP-Growth (Log Scale).

1.0.2 Analysis of Observations

- **At Low Support (5% - 10%):** We observe a massive gap in performance. **Apriori** takes significantly longer (or times out), whereas **FP-Growth** remains efficient. *Reason:* At low support, many items are considered frequent. Because `webdocs` has long transactions, the search space for Apriori explodes (checking millions of candidate combinations). FP-Growth avoids this by simply following the paths that effectively exist in the tree.
- **At High Support (50% - 90%):** The performance of both algorithms converges. *Reason:* At 90% support, very few items appear frequently enough to be counted. The computational cost becomes dominated by simple file I/O, which is similar for both.

2 Task 1.2: Dataset Creation for Runtime Replication

Objective

The goal of this task was to construct a synthetic transactional dataset that replicates the specific performance characteristics observed in Task 1.1:

1. **Low Support (5%–10%)**: Apriori should be significantly slower than FP-Growth.
2. **Medium/High Support (25%–50%)**: Apriori should outperform or match FP-Growth.

Dataset Construction Strategy

To achieve these runtime trends, we developed a Python script (`generation.py`) that generates approximately 15,000 transactions and 1000 items. The generation logic relied on two critical statistical properties:

- **High Transaction Density (Long Patterns)**: The primary bottleneck of Apriori is the combinatorial explosion of candidate generation. If a transaction has length L , it contains $2^L - 1$ potential subsets. To force Apriori to struggle at low support, we sampled transaction lengths from a heavy-tailed distribution ranging from **60 to 450 items**. These "dense" transactions create massive candidate sets (C_2, C_3) when the support threshold is low (5%), causing Apriori's runtime to spike exponentially.
- **Zipfian Item Distribution ($\alpha = 1.35$)**: We distributed item frequencies using a Zipfian distribution with exponent $\alpha = 1.35$.

$$P(\text{rank}) \propto \frac{1}{\text{rank}^{1.35}}$$

This ensures that while a few items are very frequent, the "tail" of the distribution is thin.

Experimental Results

We executed both algorithms on the constructed dataset. The runtime trends are visualized below.

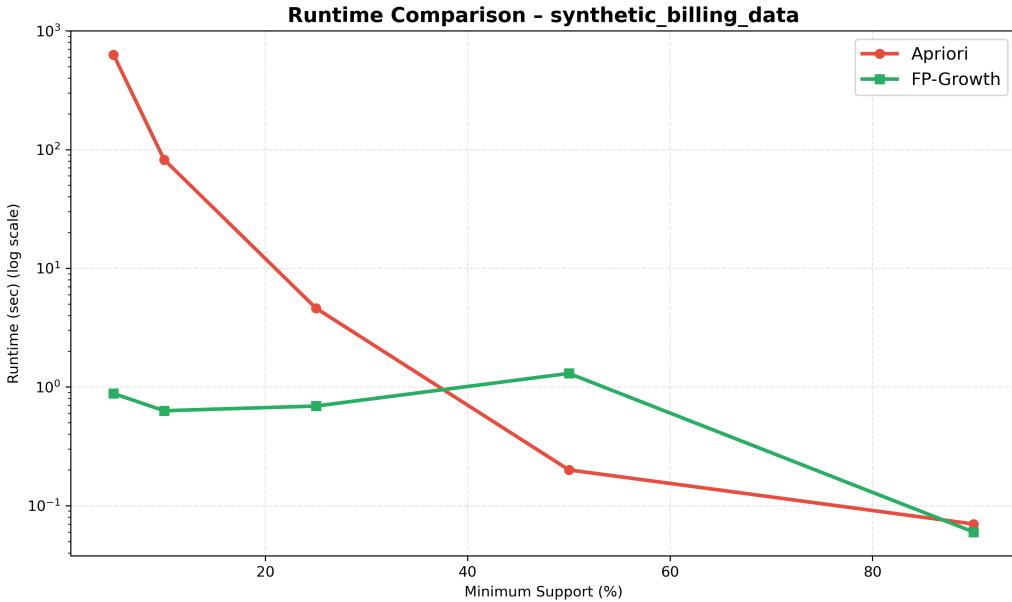


Figure 2: Runtime comparison on the Synthetic Dataset (Task 1.2).

Analysis

The resulting plot confirms that our dataset construction strategy was successful:

1. **At 5%–10% Support:** Apriori has a massive performance degradation, indicated by the steep curve/spike. This confirms that the long transaction lengths effectively overwhelmed the candidate generation step. In contrast, FP-Growth handled the dense data efficiently due to its compact tree structure.
2. **At 25%–50% Support:** A clear crossover is observed. As the support threshold eliminates the "long tail" of infrequent items (due to the Zipfian skew), Apriori's runtime drops below that of FP-Growth.

Disclaimer

Large Language Models were used during this assignment for clarifying the logic of Frequent Itemset Mining algorithms and for assistance with L^AT_EX syntax. The implementation logic, data generation strategy, and experimental conclusions are our own work.

References

- [1] R. Agrawal and R. Srikant, "Fast Algorithms for Mining Association Rules in Large Databases," *VLDB*, 1994.
- [2] J. Han, J. Pei, and Y. Yin, "Mining Frequent Patterns without Candidate Generation," *SIGMOD*, 2000.