

COL761: Assignment 1 - Part 3

Abhay Sharma
2012CS50272

Akash Dabaniya
2021CS10111

Mridul Singh
2018CS50412

February 4, 2026

1 Task 3: Graph Indexing

Problem Overview

The objective of this task is to perform efficient subgraph isomorphism queries on a large graph database. A naïve approach (checking isomorphism against every graph) is computationally prohibitive ($O(|D| \cdot NP)$).

To solve this, we implemented a **feature-based indexing strategy**[3]. We mine discriminative subgraph fragments to create binary feature vectors for the database graphs. At query time, we filter out non-promising candidates using the inclusion-exclusion principle, significantly reducing the candidate set size $|C_q|$ passed to the expensive verification step.

2 Methodology

2.1 Step 1: Discriminative Subgraph Mining

Instead of using random subgraphs or simply the most frequent ones, we developed a pipeline to identify features that maximize information gain.

1. Mining with Gaston

[2]We utilized the Gaston algorithm for Frequent Subgraph Mining (FSM).

- **Why** Based on our experiments in Part 2, Gaston consistently outperformed gSpan and FSG in runtime.
- **Support Threshold:** We set a low minimum support (approx. 2%) to capture a wide variety of potential features, ensuring we didn't miss rare but discriminative patterns.

2. Feature Selection Metric (Variance Maximization)

A feature is only useful for indexing if it can distinguish between graphs. A feature present in 100% of graphs or 0% of graphs provides zero pruning power.

We scored each mined subgraph f based on the variance of a Bernoulli trial:

$$Score(f) = P(f) \times (1 - P(f))$$

Where $P(f) = \frac{\text{support}(f)}{|D|}$.

This metric favors subgraphs with support close to **50%**. These features theoretically split the database in half, maximizing the pruning power of each bit in the index.

3. Redundancy Pruning via Jaccard Similarity

Simply picking top-scoring features often results in selecting variations of the same substructure. To ensure orthogonality in our feature space, we implemented a redundancy check:

1. For every new candidate feature, we compared its transaction ID set (T_{new}) with the transaction ID sets of already selected features ($T_{selected}$).
2. We calculated the Jaccard Similarity:

$$J(T_{new}, T_{selected}) = \frac{|T_{new} \cap T_{selected}|}{|T_{new} \cup T_{selected}|}$$

3. If $J > 0.90$, the new feature was deemed redundant (co-occurring too frequently with an existing feature) and was discarded.

2.2 Step 2: Index Construction

For the database D and query q , we constructed binary feature vectors V_D and V_q of length $K = 50$ (where K is the number of selected features).

$$V[i] = \begin{cases} 1 & \text{if feature } f_i \subseteq G \\ 0 & \text{otherwise} \end{cases}$$

This mapping was performed using ‘NetworkX’ for subgraph isomorphism checks.

2.3 Step 3: Candidate Generation (Filtering)

We generated the candidate set C_q using the necessary condition for subgraph isomorphism:

If query q is a subgraph of g , then every feature present in q must also be present in g .

Using bitwise operations, we retained a graph g_i in the candidate set only if:

$$(V_{g_i} \text{ AND } V_q) == V_q$$

This operation is extremely fast and effectively prunes the search space.

3 Conclusion

By combining the speed of Gaston for mining, a variance-based scoring metric for feature selection, and Jaccard-based pruning for diversity, our approach constructs a compact and highly discriminative index. This minimizes the size of the candidate set $|C_q|$, thereby optimizing the competitive score defined as $S_q = |R_q|/|C_q|$.

Disclaimer

Large Language Models were used during this assignment for clarifying the workflow of graph indexing pipelines and for assistance with L^AT_EX syntax. The implementation logic and experimental conclusions are our own work.

References

- [1] X. Yan and J. Han, "gSpan: Graph-Based Substructure Pattern Mining," *ICDM*, 2002.
- [2] S. Nijssen and J. Kok, "The Gaston Tool for Frequent Subgraph Mining," *Electronic Notes in Theoretical Computer Science*, 2005.
- [3] X. Yan, P. S. Yu, and J. Han, "Graph Indexing: A Frequent Structure-based Approach," *SIGMOD*, 2004.